



Healthcare Data Science Research Proposal

Machine Learning-Based Prediction of Drug Candidate Activity Using Public Bioactivity Data

*A Healthcare Data Science Approach for Early-Stage Drug
Discovery*

SUBMITTED BY

PRAVEEN R

BS-MS Student

Indian Institute of Science Education and Research (IISER) Tirupati

SUBMITTED TO

YuvaIntern

Healthcare Data Science Researcher Internship

WEEK 1 SUBMISSION

Research Project Proposal and Strategy Development

PROPOSED TOOLS & TECHNOLOGIES

Python | Pandas | NumPy | Scikit-learn | Matplotlib | Seaborn

DATE OF SUBMISSION

29/07/2026

Table of Contents

S.No.	Section	Page No.
1	Introduction to the Healthcare Issue	2
2	Literature Review Summary	4
3	Research Objectives	6
4	Research Questions	7
5	Hypothesis Formulation	8
6	Proposed Data Science Techniques	9
7	Research Methodology	12
8	Planned Python Libraries	16
9	Project Timeline	17
10	Potential Challenges and Mitigation Strategies	18
11	Alignment with Healthcare Data Science Principles	20
12	References	21

Abstract

Drug discovery is a complex, time-consuming, and expensive process that requires the evaluation of thousands of chemical compounds before identifying potential therapeutic candidates. Recent advances in healthcare data science, artificial intelligence, and machine learning have provided new opportunities to accelerate this process through bioactivity prediction using large-scale public bioactivity dataset. Such computational approaches can reduce both the cost and time associated with early-stage drug discovery.

This research proposal aims to develop a machine learning framework for predicting the biological activity of drug candidates using publicly available bioactivity datasets. The study will employ healthcare data science techniques, including data preprocessing, exploratory data analysis, feature engineering, and supervised machine learning, to classify compounds as biologically active or inactive. Python libraries such as Pandas, NumPy, Scikit-learn, Matplotlib, and Seaborn will be used to implement the proposed workflow and evaluate model performance using standard classification metrics.

The expected outcome is a systematic predictive modelling workflow that demonstrates the application of machine learning and artificial intelligence in medicinal chemistry and early-stage drug discovery. The project also aims to illustrate how healthcare data science and bioactivity prediction can support evidence-based decision-making by identifying promising drug candidates for further experimental investigation.

Keywords: Healthcare Data Science; Machine Learning; Drug Discovery; Bioactivity Prediction; Medicinal Chemistry; Artificial Intelligence; Predictive Modelling.

1. Introduction to the Healthcare Issue

Drug discovery is a fundamental component of modern healthcare, enabling the development of therapeutic agents for the prevention, management, and treatment of human diseases. The discovery of new drugs has significantly improved global health by providing effective treatments for infectious diseases, chronic illnesses, genetic disorders, and cancer. However, despite remarkable advances in biomedical science, the traditional drug discovery process remains one of the most challenging and resource-intensive areas of pharmaceutical research. The identification, optimization, and validation of a single drug candidate typically require extensive laboratory experimentation, multidisciplinary collaboration, and years of research before a compound reaches clinical use.

Conventional drug discovery follows a systematic pipeline that includes target identification, hit discovery, lead optimization, preclinical testing, clinical trials, and regulatory approval. During the early stages of this process, researchers may screen thousands to millions of chemical compounds to identify molecules with the desired biological activity against a specific therapeutic target. The majority of these compounds fail due to inadequate efficacy, toxicity, poor pharmacokinetic properties, or lack of selectivity. Consequently, the overall success rate of drug development is relatively low, while the associated financial investment and development time remain exceptionally high. These limitations have created an urgent need for more efficient and data-driven approaches capable of improving decision-making during the initial phases of drug discovery.

The rapid growth of biomedical databases, computational chemistry, and artificial intelligence has transformed the landscape of pharmaceutical research. Publicly accessible resources such as ChEMBL, PubChem, DrugBank, UniProt, and the Protein Data Bank provide large volumes of curated chemical, biological, and pharmacological information that can be utilized for computational analysis. The availability of these datasets has enabled researchers to apply healthcare data science techniques for

extracting meaningful patterns from complex biological data. By integrating computational methods with pharmaceutical research, scientists can prioritize promising compounds before experimental validation, thereby reducing unnecessary laboratory work and accelerating the discovery pipeline.

Among the various computational approaches, machine learning has emerged as one of the most promising technologies for drug discovery. Machine learning algorithms can identify complex relationships between molecular features and biological activity by learning from previously characterized compounds. Instead of relying solely on experimental screening, predictive models can estimate the likelihood that a chemical compound will exhibit biological activity against a particular target. This process, commonly referred to as bioactivity prediction, allows researchers to rapidly evaluate large collections of compounds and prioritize candidates with higher therapeutic potential. Such predictive capabilities contribute to improved resource utilization and support more informed decision-making during early-stage pharmaceutical research.

Healthcare data science combines principles from statistics, computer science, artificial intelligence, and domain-specific biomedical knowledge to solve healthcare-related problems using data-driven methodologies. Within drug discovery, healthcare data science enables the integration of chemical information, biological assays, molecular descriptors, and computational algorithms to develop predictive models capable of supporting pharmaceutical research. The combination of healthcare data science and machine learning has demonstrated considerable potential in areas such as virtual screening, drug repurposing, toxicity prediction, target identification, and lead optimization, making it an increasingly valuable component of modern medicinal chemistry.

The present research proposal focuses on the development of a machine learning-based predictive modelling framework for bioactivity prediction using publicly available bioactivity datasets. The proposed study aims to investigate whether supervised machine learning techniques can accurately classify chemical compounds as biologically active or inactive based on molecular and bioactivity information. By employing standard

healthcare data science methodologies—including data preprocessing, exploratory data analysis, feature engineering, predictive modelling, and model evaluation—the project seeks to establish a systematic workflow for computational drug discovery. The findings of this study are expected to demonstrate the potential of machine learning-based predictive modelling to complement conventional pharmaceutical research and support more efficient identification of promising drug candidates for subsequent experimental validation.

2. Literature Review Summary

Drug discovery has traditionally relied on extensive experimental screening of chemical compounds to identify molecules with potential therapeutic activity. However, this approach is associated with high costs, lengthy development timelines, and low success rates, creating a significant productivity challenge within the pharmaceutical industry (Paul et al., 2010). The increasing availability of publicly accessible chemical and biological data has encouraged researchers to explore computational approaches that can improve the efficiency of early-stage drug discovery. Consequently, machine learning and healthcare data science have become important components of modern pharmaceutical research.

Large public databases such as ChEMBL, PubChem, DrugBank, and the Protein Data Bank (PDB) provide comprehensive information on chemical structures, biological activities, molecular targets, and pharmacological properties. Among these resources, ChEMBL has become one of the most widely used databases for machine learning applications because it contains manually curated bioactivity data extracted from scientific literature (Gaulton et al., 2017). These databases provide the foundation for developing predictive models capable of identifying biologically active compounds before laboratory experimentation.

Recent studies have demonstrated that machine learning algorithms can effectively predict the biological activity of chemical compounds using molecular descriptors and bioactivity datasets (Vamathevan et al., 2019). Supervised learning techniques such as Random Forest, Support Vector Machine, Gradient Boosting, and Artificial Neural Networks have shown

promising performance in classifying compounds as active or inactive against specific biological targets. More recently, deep learning approaches, including Graph Neural Networks and transformer-based molecular models, have further improved predictive accuracy by learning complex structural relationships directly from molecular representations.

Researchers have also emphasized the importance of proper data preprocessing, feature engineering, model validation, and performance evaluation in developing reliable predictive models. Techniques such as molecular fingerprint generation, descriptor calculation, feature selection, and cross-validation have been shown to improve model robustness and reduce overfitting. Performance is commonly assessed using evaluation metrics including accuracy, precision, recall, F1-score, Receiver Operating Characteristic–Area Under the Curve (ROC–AUC), and Matthews Correlation Coefficient (MCC), depending on the characteristics of the dataset.

Despite these advancements, several challenges remain. Public bioactivity datasets often contain missing values, class imbalance, experimental variability, and inconsistent annotations, which can negatively affect model performance. Furthermore, many published studies focus primarily on maximizing prediction accuracy without providing easily reproducible workflows suitable for students and early-stage researchers. There remains a need for systematic, transparent, and reproducible machine learning pipelines that integrate healthcare data science principles with publicly available bioactivity datasets for educational and research applications.

The present research proposal addresses this gap by proposing a structured machine learning workflow for bioactivity prediction using publicly available datasets. The proposed framework incorporates data acquisition, preprocessing, exploratory data analysis, feature engineering, supervised machine learning, model evaluation, and result interpretation using Python. By emphasizing reproducibility, transparency, and practical implementation, this study aims to demonstrate how healthcare data science and machine learning can support early-stage drug discovery while providing a foundation for future computational medicinal chemistry research.

3. Research Objectives

3.1 Primary Objective

To develop a machine learning-based predictive modelling framework for identifying the biological activity of chemical compounds using publicly available bioactivity datasets, thereby demonstrating the application of healthcare data science and artificial intelligence in early-stage drug discovery.

3.2 Specific Objectives

1. To investigate the role of machine learning and healthcare data science in accelerating the early stages of drug discovery.
2. To acquire and utilize publicly available bioactivity datasets from databases such as ChEMBL and PubChem for predictive modelling.
3. To preprocess and prepare the collected dataset by handling missing values, removing inconsistencies, and performing appropriate feature engineering techniques.
4. To perform exploratory data analysis (EDA) to understand the characteristics, distribution, and relationships within the bioactivity data.
5. To develop supervised machine learning models capable of classifying chemical compounds as biologically active or biologically inactive.
6. To evaluate the predictive performance of the developed models using standard machine learning evaluation metrics, including Accuracy, Precision, Recall, F1-score, and Receiver Operating Characteristic–Area Under the Curve (ROC–AUC).
7. To establish a transparent, systematic, and reproducible machine learning workflow using Python and relevant healthcare data science libraries.
8. To demonstrate the potential of computational approaches in supporting medicinal chemistry research by prioritizing promising drug candidates for further experimental validation.

4. Research Questions

This research proposal seeks to address the following research questions:

1. Can machine learning models accurately predict the biological activity of chemical compounds using publicly available bioactivity datasets?
2. How do data preprocessing and feature engineering techniques influence the predictive performance of machine learning models for bioactivity prediction?
3. How effectively can publicly available bioactivity databases, such as ChEMBL and PubChem, support machine learning-based drug discovery research?
4. Which supervised machine learning algorithm provides the most reliable performance for classifying chemical compounds as biologically active or biologically inactive?
5. How effectively can the predictive performance of machine learning models be evaluated using metrics such as Accuracy, Precision, Recall, F1-score, and Receiver Operating Characteristic–Area Under the Curve (ROC–AUC)?
6. Can a transparent, systematic, and reproducible machine learning workflow be developed using Python to support healthcare data science applications in early-stage drug discovery?
7. How can machine learning-based bioactivity prediction assist researchers in prioritizing promising drug candidates for further experimental validation?
8. To what extent can healthcare data science and artificial intelligence contribute to improving the efficiency of computational drug discovery?

5. Hypothesis Formulation

A hypothesis is a testable statement that predicts the expected outcome of a research study. Based on the research objectives, research questions, and existing literature, the following hypotheses are proposed.

5.1 Null Hypothesis (H_0)

Machine learning models developed using publicly available bioactivity datasets do not achieve statistically significant predictive performance in classifying chemical compounds as biologically active or biologically inactive.

5.2 Alternative Hypothesis (H_1)

Machine learning models developed using publicly available bioactivity datasets achieve statistically significant predictive performance in classifying chemical compounds as biologically active or biologically inactive, thereby demonstrating their potential application in early-stage drug discovery.

5.3 Expected Outcome

Based on findings reported in previous studies and recent advances in healthcare data science, it is anticipated that supervised machine learning models trained on well-preprocessed public bioactivity datasets will demonstrate reliable predictive performance. The proposed workflow is expected to establish a transparent, systematic, and reproducible computational framework capable of supporting bioactivity prediction, prioritizing promising drug candidates, and assisting early-stage drug discovery through data-driven decision-making.

6. Proposed Data Science Techniques

The proposed study will employ a systematic combination of healthcare data science and machine learning techniques to develop a predictive framework for bioactivity prediction using publicly available bioactivity datasets. The selected techniques have been chosen based on their widespread application in computational drug discovery, medicinal chemistry, and predictive modelling research. The proposed workflow aims to establish a transparent, reproducible, and efficient computational approach for identifying biologically active chemical compounds.

6.1 Data Collection

Publicly available bioactivity datasets will be obtained from established biomedical databases such as ChEMBL and PubChem. These databases contain experimentally validated information on chemical compounds, molecular structures, canonical SMILES, biological targets, bioactivity measurements, and molecular descriptors, making them suitable for machine learning-based bioactivity prediction.

6.2 Data Preprocessing

The collected dataset will undergo preprocessing to improve data quality and ensure compatibility with machine learning algorithms. Planned preprocessing activities include:

- Removal of duplicate records.
- Handling missing or incomplete values.
- Data normalization and standardization where appropriate.
- Feature selection and feature engineering.
- Preparation of the dataset for supervised learning.

6.3 Exploratory Data Analysis (EDA)

Exploratory Data Analysis will be performed to understand the characteristics of the dataset before model development. This stage will include:

- Statistical summary of variables.
- Distribution analysis.
- Correlation analysis.
- Identification of class imbalance.
- Detection of outliers and anomalies.
- Data visualization using graphs and charts.

EDA will assist in identifying meaningful patterns and preparing the dataset for predictive modelling.

6.4 Machine Learning Techniques

Supervised machine learning algorithms will be investigated for bioactivity prediction. Depending on dataset characteristics, the following algorithms may be considered:

- Logistic Regression
- Decision Tree
- Random Forest
- Support Vector Machine (SVM)
- Gradient Boosting
- Artificial Neural Networks (ANN)

The performance of these algorithms will be compared to determine the most suitable model for predicting the biological activity of chemical compounds as biologically active or biologically inactive.

6.5 Model Evaluation

The predictive models will be evaluated using standard classification metrics to ensure reliability and robustness. The evaluation metrics will include:

- Accuracy
- Precision
- Recall
- F1-score
- Receiver Operating Characteristic–Area Under the Curve (ROC–AUC)

Where appropriate, cross-validation techniques will also be employed to assess model generalizability and minimize the risk of overfitting. These evaluation metrics will provide a comprehensive assessment of model performance and predictive capability.

6.6 Data Visualization

Data visualization techniques will be employed throughout the study to communicate analytical findings effectively. Planned visualizations include:

- Histograms
- Bar charts
- Scatter plots
- Box plots
- Correlation heatmaps
- Pair plots (where applicable)
- Receiver Operating Characteristic (ROC) curves

These visualizations will facilitate interpretation of the dataset, identify meaningful trends and relationships, and support the evaluation of predictive model performance.

6.7 Reproducible Workflow

The entire workflow will be implemented using Python to ensure transparency, reproducibility, and scalability. All stages of the study—including data acquisition, preprocessing, exploratory analysis, feature engineering, model development, evaluation, and visualization—will be systematically documented. This structured workflow will facilitate reproducibility and support future research in computational drug discovery, medicinal chemistry, and healthcare data science.

7. Research Methodology

7.1 Research Design

This study will adopt a quantitative and computational research approach based on healthcare data science and machine learning methodologies. The research will focus on developing a predictive modelling framework capable of classifying chemical compounds as biologically active or biologically inactive using publicly available bioactivity datasets. The overall methodology will follow a systematic workflow consisting of data collection, preprocessing, exploratory data analysis, feature engineering, model development, performance evaluation, result interpretation, and documentation.

7.2 Data Source

The primary datasets will be obtained from publicly available biomedical databases such as ChEMBL and PubChem. These databases provide experimentally validated information on chemical compounds, molecular structures, molecular descriptors, canonical SMILES, biological targets, and bioactivity measurements. The selected datasets will serve as the foundation for developing and evaluating machine learning models for bioactivity prediction. The dataset selected for this study will be chosen based on its relevance, completeness, and suitability for supervised machine learning-based bioactivity prediction.

7.3 Data Preprocessing

Prior to model development, the collected dataset will undergo preprocessing to improve data quality and consistency. The preprocessing stage will include:

- Removal of duplicate entries.
- Handling of missing or incomplete values.
- Data normalization and standardization where appropriate.
- Encoding of categorical variables, if required.
- Removal of highly correlated or redundant features, where appropriate.
- Feature selection and feature engineering.
- Preparation of the dataset for supervised machine learning.

These preprocessing procedures will improve dataset reliability and reduce the influence of noise and inconsistencies on predictive performance.

7.4 Exploratory Data Analysis

Exploratory Data Analysis (EDA) will be performed to understand the characteristics of the collected bioactivity dataset before predictive model development. This stage will assist in identifying data quality issues, understanding feature distributions, detecting relationships among variables, and supporting informed decision-making during model development.

The exploratory analysis will include:

- Statistical summary of variables.
- Analysis of data distributions.
- Identification of missing values and outliers.
- Correlation analysis between molecular descriptors and bioactivity variables.
- Detection of class imbalance.

- Visualization of data using appropriate graphs and charts.

Visualization techniques such as histograms, box plots, scatter plots, and correlation heatmaps will be employed to identify patterns and relationships within the dataset. The insights obtained during EDA will guide feature engineering and improve the effectiveness of subsequent predictive modelling.

7.5 Model Development

Supervised machine learning algorithms will be implemented using Python to predict the biological activity of chemical compounds. Multiple classification algorithms will be investigated, including Logistic Regression, Decision Tree, Random Forest, Support Vector Machine (SVM), Gradient Boosting, and Artificial Neural Networks (ANN).

The dataset will be divided into training and testing subsets using an appropriate train-test split strategy (e.g., 80:20). Each model will be trained using the training dataset and evaluated using the testing dataset. The comparative performance of all models will be analysed to identify the most suitable algorithm for bioactivity prediction.

7.6 Model Evaluation

The developed machine learning models will be evaluated using widely accepted classification performance metrics to ensure reliability and predictive capability. The evaluation metrics will include:

- Accuracy
- Precision
- Recall
- F1-score
- Receiver Operating Characteristic–Area Under the Curve (ROC–AUC)

Where appropriate, cross-validation techniques will also be employed to assess model robustness, reduce overfitting, and evaluate model generalizability. The best-performing model will be selected based on its overall predictive performance across these evaluation metrics.

7.7 Software and Programming Environment

The computational workflow will be implemented using Python within the Jupyter Notebook environment. The primary Python libraries proposed for this study include:

- Pandas
- NumPy
- Scikit-learn
- Matplotlib
- Seaborn

These libraries will be utilized for data preprocessing, exploratory data analysis, feature engineering, machine learning model development, visualization, and performance evaluation.

7.8 Ethical Considerations

This research will utilize only publicly available and non-identifiable bioactivity datasets obtained from recognized biomedical databases. No confidential patient information, clinical records, or personally identifiable data will be collected or analysed. Consequently, the proposed study does not involve direct human participation and therefore presents minimal ethical concerns. The research will adhere to the principles of responsible data usage, scientific integrity, transparency, reproducibility, and ethical application of artificial intelligence in healthcare research.

7.9 Methodological Workflow

The proposed research methodology will follow the workflow illustrated below:

Public Bioactivity Data Collection (ChEMBL/PubChem)

↓

Data Cleaning and Preprocessing

↓

Exploratory Data Analysis (EDA)

↓

Feature Engineering and Feature Selection

↓

Machine Learning Model Development

↓

Model Training and Testing

↓

Performance Evaluation

↓

Prediction of Bioactivity

↓

Interpretation of Results and Documentation

8. Planned Python Libraries

The proposed research will be implemented using Python and its scientific computing ecosystem. The following libraries have been selected based on their capabilities in data preprocessing, exploratory data analysis, machine learning, and data visualization. These libraries are widely used in healthcare data science, computational drug discovery, and bioinformatics research.

Python Library	Purpose in the Proposed Research
Pandas	Data loading, cleaning, preprocessing, manipulation, and analysis of bioactivity datasets.
NumPy	Numerical computations, matrix operations, and mathematical calculations required during data preprocessing and model development.
Scikit-learn	Implementation of supervised machine learning algorithms, feature engineering, model training, testing, and performance evaluation.
Matplotlib	Generation of scientific graphs and visualizations, including histograms, scatter plots, and ROC curves.

Seaborn	Statistical data visualization, including correlation heatmaps, box plots, pair plots, and distribution analysis.
Jupyter Notebook	Interactive Python development environment for implementing, documenting, and reproducing the complete computational workflow.

The selected Python libraries collectively provide a comprehensive framework for implementing the proposed machine learning workflow, ensuring reproducibility, computational efficiency, and reliable analysis of public bioactivity datasets. Additional libraries may be incorporated if required during the implementation phase, depending on the characteristics of the selected dataset and the performance of the developed machine learning models.

9. Project Timeline

The proposed research project will be conducted over the four-week internship period. Each week is designed to build progressively upon the previous stage, ensuring the systematic development, evaluation, and interpretation of the machine learning framework for bioactivity prediction.

Internship Week	Planned Activities	Expected Deliverables
Week 1	Literature review, identification of the healthcare challenge, formulation of research objectives, research questions, hypothesis, proposed data science techniques, and research methodology.	Research proposal document outlining the complete project plan.

Week 2	Collection of public bioactivity datasets, data preprocessing, data cleaning, exploratory data analysis (EDA), statistical analysis, and visualization of dataset characteristics.	Documentation of dataset preparation, exploratory data analysis, Python code snippets, visualizations, and preliminary findings.
Week 3	Development of supervised machine learning models, model training, validation, testing, performance evaluation, and comparative analysis of different algorithms.	Documentation of predictive model development, evaluation metrics, graphical results, and interpretation of model performance.
Week 4	Comprehensive data visualization, interpretation of analytical findings, preparation of recommendations, compilation of the final report, and project documentation.	Final project report containing visualizations, conclusions, recommendations, and complete documentation of the computational workflow.

\

10. Potential Challenges and Mitigation Strategies

Although the proposed research is computational in nature, several challenges may arise during different stages of the study. Appropriate mitigation strategies have been identified to minimize their impact and ensure the successful completion of the project.

Potential Challenge	Possible Impact	Mitigation Strategy
Missing or incomplete bioactivity data	Reduced data quality and model reliability	Apply appropriate data cleaning techniques, handle missing values systematically, and select high-quality publicly available datasets.
Class imbalance between active and inactive compounds	Biased model predictions and reduced classification performance	Use appropriate sampling techniques and evaluate models using multiple performance metrics such as Precision, Recall, F1-score, and ROC-AUC.
High dimensionality of molecular descriptors	Increased computational complexity and potential overfitting	Apply feature selection and feature engineering techniques to retain only the most informative features.
Model overfitting	Poor generalization to unseen data	Employ train-test splitting, cross-validation, and appropriate model parameter tuning to improve model robustness.
Computational limitations	Increased execution time during model training	Optimize the computational workflow by selecting efficient machine learning algorithms, reducing unnecessary features through feature selection, and using optimized

		Python libraries to improve execution efficiency.
Data inconsistencies from multiple public databases	Reduced predictive accuracy and reduced reproducibility	Standardize data formats and perform thorough preprocessing before model development.
Limited project duration	Restricted time for extensive model optimization	Prioritize the implementation of a systematic workflow and focus on developing reliable baseline predictive models within the available timeframe.

11. Alignment with Healthcare Data Science Principles

The proposed research closely aligns with the fundamental principles of healthcare data science by integrating biomedical data, computational techniques, and machine learning to address a significant challenge in early-stage drug discovery. By utilizing publicly available bioactivity datasets, the study aims to develop predictive models capable of identifying biologically active chemical compounds, thereby supporting evidence-based decision-making in pharmaceutical research.

The proposed framework follows a comprehensive healthcare data science workflow encompassing data acquisition, preprocessing, exploratory data analysis, feature engineering, predictive modelling, model evaluation, and interpretation of results. This systematic and reproducible approach demonstrates the practical application of healthcare data science methodologies in computational drug discovery and medicinal chemistry.

The integration of machine learning techniques with large-scale bioactivity datasets has the potential to improve the efficiency of identifying promising drug candidates while reducing the time, cost, and experimental effort associated with conventional drug discovery approaches. Consequently, the proposed methodology supports data-driven prioritization of compounds for further experimental validation and contributes to more efficient pharmaceutical research.

The study also adheres to the principles of responsible healthcare data science by utilizing publicly available, non-identifiable datasets obtained from recognized biomedical databases. The proposed workflow emphasizes transparency, reproducibility, scientific integrity, and the ethical application of artificial intelligence in healthcare research.

Overall, this research demonstrates how healthcare data science can facilitate computational medicinal chemistry by transforming complex biomedical data into meaningful predictive insights. The proposed framework highlights the growing role of machine learning and artificial intelligence in supporting innovative, efficient, and evidence-based approaches to modern drug discovery.

12. References

Berman, H. M., Westbrook, J., Feng, Z., Gilliland, G., Bhat, T. N., Weissig, H., Shindyalov, I. N., & Bourne, P. E. (2000). *The Protein Data Bank*. *Nucleic Acids Research*, 28(1), 235–242. <https://doi.org/10.1093/nar/28.1.235>

Gaulton, A., Hersey, A., Nowotka, M., Bento, A. P., Chambers, J., Mendez, D., Mutowo, P., Atkinson, F., Bellis, L. J., Cibrián-Uhalte, E., Davies, M., Dedman, N., Karlsson, A., Magariños, M. P., Overington, J. P., Papadatos, G., Smit, I., & Leach, A. R. (2017). *The ChEMBL database in 2017*. *Nucleic Acids Research*, 45(D1), D945–D954. <https://doi.org/10.1093/nar/gkw1074>

Kim, S., Chen, J., Cheng, T., Gindulyte, A., He, J., He, S., Li, Q., Shoemaker, B. A., Thiessen, P. A., Yu, B., Zaslavsky, L., Zhang, J., & Bolton, E. E. (2023). *PubChem 2023 update*. *Nucleic Acids Research*, 51(D1), D1373–D1380. <https://doi.org/10.1093/nar/gkac956>

Paul, S. M., Mytelka, D. S., Dunwiddie, C. T., Persinger, C. C., Munos, B. H., Lindborg, S. R., & Schacht, A. L. (2010). *How to improve R&D productivity: The pharmaceutical industry's grand challenge*. *Nature Reviews Drug Discovery*, 9(3), 203–214. <https://doi.org/10.1038/nrd3078>

Vamathevan, J., Clark, D., Czodrowski, P., Dunham, I., Ferran, E., Lee, G., Li, B., Madabhushi, A., Shah, P., Spitzer, M., & Zhao, S. (2019). *Applications of machine learning in drug discovery and development*. *Nature Reviews Drug Discovery*, 18(6), 463–477. <https://doi.org/10.1038/s41573-019-0024-5>

Wishart, D. S., Feunang, Y. D., Guo, A. C., Lo, E. J., Marcu, A., Grant, J. R., Sajed, T., Johnson, D., Li, C., Sayeeda, Z., Assempour, N., Iynkkaran, I., Liu, Y., Maciejewski, A., Gale, N., Wilson, A., Chin, L., Cummings, R., Le, D., ... Wilson, M. (2018). *DrugBank 5.0: A major update to the DrugBank database for 2018*. *Nucleic Acids Research*, 46(D1), D1074–D1082. <https://doi.org/10.1093/nar/gkx1037>

